

# NGUYEN VAN BINH

(+65) 83073859 | [nguyen.binh@u.nus.edu](mailto:nguyen.binh@u.nus.edu) | [linkedin.com/in/nvbinh/](https://www.linkedin.com/in/nvbinh/) | [github.com/nvbinh15](https://github.com/nvbinh15) | [nvbinh.com](https://nvbinh.com) | [medium.com/@nvbinh](https://medium.com/@nvbinh)

## SKILLS

<b>Programming Languages</b>	Python, C/C++, Java, RISC Assembly
<b>AI/ML Domains</b>	Computer Vision, Natural Language Processing, Graph ML, MLOps
<b>AI/ML Tools</b>	Scikit-Learn, PyTorch, TensorFlow, Keras, AWS SageMaker, MLflow
<b>Database/Data Processing</b>	RDBMS, PostgreSQL, Hadoop MapReduce, Spark, DVC
<b>Web Development</b>	FastAPI, Django, Django REST Framework, HTML/CSS/JavaScript
<b>Others</b>	Docker, Cloud Platforms (Amazon Web Services, Azure), Computer Networks, Airflow, Linux

## EDUCATION

**National University of Singapore** **Aug 2020 – now (Expected: May 2024)**

Bachelor of Engineering, Computer Engineering, Minors in Management and Data Engineering

- Grade Point Average (GPA): 4.91/5.0 – Dean's List in all eligible semesters
- NUS Overseas Colleges Program, NOC Paris (2023): participated in entrepreneurship program at Université PSL

## EXPERIENCE

**AI & Computer Vision Engineer Intern, Vebits (Singapore)** **Jan 2024 – now**

- Train and quantize ResNet-50 and YOLOv7 with ByteTrack for image classification, multi-object detection and tracking on Rockchip RK3588 board, maximize NPU utilization, and run inference in real-time at 30fps
- Capture and stream video from camera on RK3588 to RTSP server at 1080p resolution and integrate to ML pipeline doing on chip inference with maximum 500ms latency by enabling MPP hardware acceleration codec

**Data Science Intern, Science Feedback (Paris, France)** **Feb 2023 – Aug 2023**

- Developed web domain credibility ranking models with graph neural networks using PyTorch Geometric and NetworkX, built end-to-end ML pipeline, and deployed models in production (making prediction for > 1 million domains) using FastAPI
- Initiated and established data infrastructure (data warehouse, computation, orchestration) on Azure Cloud Services, enhancing data security, quality, and integrity, improving efficiency of data and software teams, and fastening ML life cycle by 40%
- Evaluated and customised social media data scraping tools to provide insights for scientists doing fact-checking task

**Research Assistant, SEEDER Group – National University of Singapore** **Jan 2022 – Dec 2022**

- Analysed tradeoffs in learning algorithms for spiking fully connected and convolutional neural networks
- Framed architecture and wrote assembly instructions for spiking neural network hardware accelerators on FPGA

**Teaching Assistant, School of Computing – National University of Singapore** **Aug 2021 – Dec 2022**

- Placed on the Honour List of Student Tutors for Excellence in Teaching with 4.8/5 feedback point (faculty average: 4.2)
- Conducted classes for Master-level IT5003 Data Structures and Algorithms (covered DSA concepts and implementation in Python) and CS2102 Database System (discussed Relational Algebra, ER model, SQL, and Database Normalization)

**Machine Learning Engineer Intern, Neuron Mobility (Singapore)** **May 2022 – Aug 2022**

- Trained and deployed ML models for scooter parking image classification on AWS SageMaker achieving more than 95% accuracy, detecting bad parking behaviours with 0.98 precision, enhancing user experience, and saving US\$200,000 annually
- Proposed and coded backend for an in-house annotation platform utilizing Django REST Framework and PostgreSQL, streamlining collaboration between ML and data annotation teams, leading to time savings of 5+ hours per week

## PROJECTS

**Quora Duplicate Question Detection and Generation** ([source code](#) | [report](#))

- Led team of 6, built duplicate question identification models using RNN, GRU, LSTM, fastText and word2vec embeddings, and fine-tuned BERT family LLMs culminating in 90.1% test accuracy rate and 0.87 F1-score
- Constructed generative language models by utilizing byte-pair encoding, n-grams, MLP, RNN, and decoder-only transformers accomplishing test 76.8 perplexity; augmented generated questions to classification task in a semi-supervised manner

**Laser Tag** ([demo](#) | [source code](#) – capstone project): Augmented Reality shooting game

- Collected data and created an ML model classifying actions based on IMU data mounted on players' hands attaining 96.14% test accuracy; translated model into high-level synthesis (C++) and ran inference on Avnet Ultra96-V2 board
- Designed architecture, optimized inter-component communication overhead through 20% data packet size reduction

## AWARDS & HACKATHON

- Champions, Vietnam National AWS DeepRacer League, Dec 2019: formulated reward functions and trained reinforcement learning models loaded to an autonomous car completing a racing lap in 9.625s

## EXTRA-CURRICULAR ACTIVITIES

- NUS School of Computing Project Mentor, 2023-2024: guiding 2 Computer Engineering freshmen, offering advice and guidance on academic and non-academic matters, internship, overseas experiences, and university life navigation